Collatz 81

September 26, 2012

The Nth *Collatz sequence* is defined by the following process, starting with the integer N:

1. if the number is even, divide it by 2,

2. if the number is odd, triple it and add 1.

Continue this process until you reach 1.

We will estimate the number of different values in the first million Collatz sequences, a stream of 132,434,271 values. To make this interesting, we pretend that we have at our disposal a Sinclair ZX81 with 1 kB of memory.

c_1: 1
c_2: 2, 1
c_3: 3, 10, 5, 16, 8, 4, 2, 1
c_4: 4, 2, 1
c_5: 5, 16, 8, 4, 2, 1
c_6: 6, 3, 10, 5, 16, 8, 4, 2, 1.
c_7: 7, 22, 11, 34, 17, 52, 26, 13, 40, 20, 10, 5, 16, 8, 4, 2, 1.
c_8: 8, 4, 2, 1
c_9: 9, 28, 14, 7, 22, 11, 34, 17, 52, 26, 13, 40, 20, 10, 5, 16, 8, 4, 2, 1.
c_{10}: 10, 5, 16, 8, 4, 2, 1.

Algorithm D

Algorithm D approximates the number of distinct elements in a stream of values $x_1, x_2, \ldots$. We use a pairwise independent hash function h mapping the stream elements to a range $\{0, \ldots, R-1\}$, where R an integer much larger than the number of distinct elements we can imagine. (Typically one will use R to be the number of positive integers available in a machine word.) After seeing i elements, the algorithm stores a set $V \subseteq \{ (h(x_j), x_j) \mid j = 1, \ldots, i \}$, corresponding to the t distinct elements having the smallest hash values. (Ties can be broken arbitrarily.) When we see the ith element x_i, we compute its hash value $h(x_i)$. If this is smaller than some hash value in V, and $(h(x_i), x_i) \notin V$, we add $(h(x_i), x_i)$ to V and throw away an element having the largest hash value. To estimate the number of distinct elements seen we look at the largest hash value v in H, and compute tR/v. The relative error of this estimate can be shown to be a factor $1 \pm O(t^{-0.5})$ with probability 3/4. To increase the confidence to arbitrarily close to 1, run the algorithm several times (in "parallel") and take the median of the estimates.

Collatz 81 Report

by Mats Rydberg and Martin Larsson[1]

1. Dictionaries

We let C_N denote the ("flattened") sequence of the sequences $c_1, \ldots, c_N$. For instance, C_3 is 1, 2, 1, 3, 10, 5, 16, 8, 4, 2, 1.

The following table gives the maximum value appearing in C_N, the number of distinct values in C_N (i.e., the cardinality of C_N viewed as a set), and the total length of the sequence $|c_1| + \cdots + |c_N|$, for increasing values of N.[2]

N	$\max C_N$	$|C_N|$	$\text{len}(C_N)$
10	52	22	77
100	9,232	251	3,242
1,000	250,504	2,228	60,542
10,000	27,117,424	21,664	859,666
100,000	1,570,824,736	217,212	10,853,840
1,000,000	56,991,483,520	2,168,611	132,434,424

2. Quadratic Time

The first solution in small space uses the following idea: For given N, produce every value of C_N (without storing the entire sequence!) to determine $\max C_N$. Start a counter at 0. Then, for every $i = 1, \ldots, \max C_N$, produce the entire sequence to see if i appears. If so, increase the counter. The running time will grow as the product of $\text{len}\, C_N$ and $\max C_n$.

The largest N for which this idea works within 60 seconds on gamma-10 was 446. On SX-box we could do $N = 702$ under 60 seconds.

3. Randomized Approximation in Small Space

The following table shows the output of our implementation of Algorithm D, together with the error (in percent) relative to the correct values of $|C_N|$ computed in the first part of the report.

Here we have set $t = 500$.

N	output	relative error
10	23	4,35%
100	252	0,40%
1,000	2,226	0,09%
10,000	21,546	0,54%
100,000	219,945	1,24%
1,000,000	2,456,243	11,7%

Perspective

The approximation algorithm is from XXX.

For a nice project idea, install a ZX81 emulator and actually solve the problem on that, preferably in ZX Basic. (If you're really cool, get a ZX81. Or make it run on the hardware of your washing machine or one of those annoying annoying greeting cards that play *Happy Birthday*.)